

```
In [25]: import pandas as pd
import numpy as np
import missingno as msno
import plotly.express as px
```

Data Cleaning

```
In [26]: df = pd.read_csv('C:/Users/19229/OneDrive/Desktop/data1030-fall2025/data1030-project.csv')
df.shape
df.head()
df.dtypes
#df = df.drop(['customerID'], axis = 1), already done
```

```
Out[26]: customerID      object
gender          object
SeniorCitizen    int64
Partner          object
Dependents       object
tenure           int64
PhoneService     object
MultipleLines    object
InternetService  object
OnlineSecurity   object
OnlineBackup     object
DeviceProtection object
TechSupport      object
StreamingTV      object
StreamingMovies  object
Contract         object
PaperlessBilling object
PaymentMethod    object
MonthlyCharges   float64
TotalCharges     object
Churn            object
dtype: object
```

```
In [27]: msno.matrix(df)
df['TotalCharges'] = pd.to_numeric(df.TotalCharges, errors='coerce')
print(df.isnull().sum())
df[np.isnan(df['TotalCharges'])]
df['TotalCharges'] = df['TotalCharges'].fillna(df['TotalCharges'].mean())
print(df.isnull().sum())

df["SeniorCitizen"] = df["SeniorCitizen"].map({0:"No", 1:"Yes"}).astype("category")
numerical_cols = ['tenure', 'MonthlyCharges', 'TotalCharges']
df[numerical_cols].describe()
```

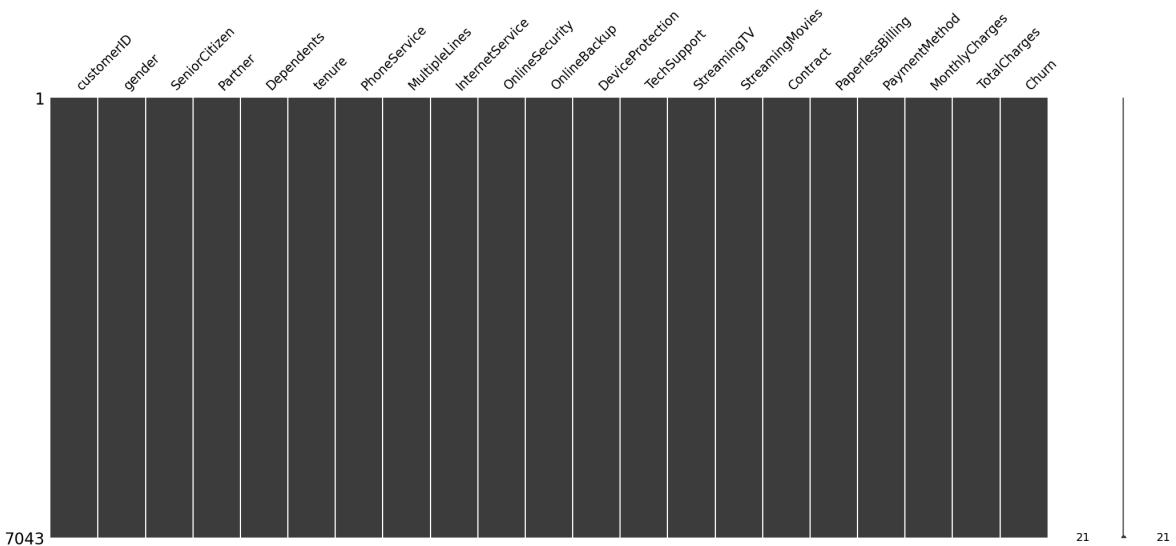
```

customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents     0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    11
Churn           0
dtype: int64
customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents     0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    0
Churn           0
dtype: int64

```

Out[27]:

	tenure	MonthlyCharges	TotalCharges
count	7043.000000	7043.000000	7043.000000
mean	32.371149	64.761692	2283.300441
std	24.559481	30.090047	2265.000258
min	0.000000	18.250000	18.800000
25%	9.000000	35.500000	402.225000
50%	29.000000	70.350000	1400.550000
75%	55.000000	89.850000	3786.600000
max	72.000000	118.750000	8684.800000



EDA

```
In [28]: import matplotlib.pyplot as plt
import os
plt.rcParams["figure.figsize"] = (7,4)

os.makedirs("reports/figures", exist_ok=True)

target = "Churn"

_churn_num = df[target].map({"No":0, "Yes":1})#temporary switch to 0/1 for matrix
```

```
In [29]: print("Shape:", df.shape)
print("Missing per column after cleaning:\n", df.isna().sum()[df.isna().sum()>0])
print("Churn distribution:\n", df[target].value_counts())
print("Positive rate:", _churn_num.mean().round(3))
print("\nNumeric summary:\n", df[numerical_cols].describe().T)
```

Shape: (7043, 21)
 Missing per column after cleaning:
 Series([], dtype: int64)
 Churn distribution:
 Churn
 No 5174
 Yes 1869
 Name: count, dtype: int64
 Positive rate: 0.265

Numeric summary:

	count	mean	std	min	25%	50%	\
tenure	7043.0	32.371149	24.559481	0.00	9.000	29.00	
MonthlyCharges	7043.0	64.761692	30.090047	18.25	35.500	70.35	
TotalCharges	7043.0	2283.300441	2265.000258	18.80	402.225	1400.55	

	75%	max
tenure	55.00	72.00
MonthlyCharges	89.85	118.75
TotalCharges	3786.60	8684.80

```
In [30]: counts = df[target].value_counts().reindex(["No","Yes"])
labels = counts.index
```

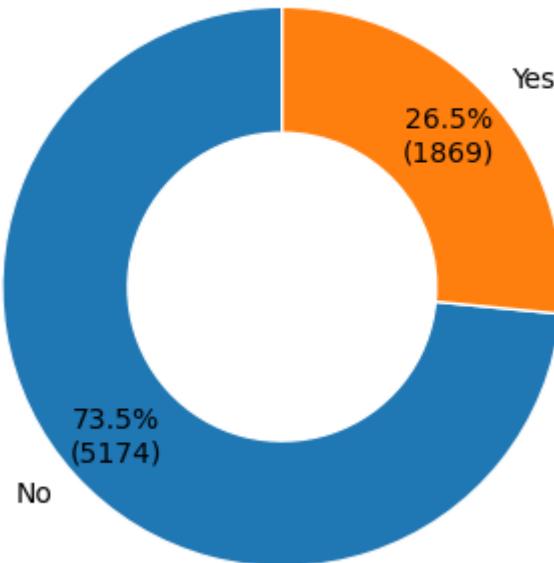
```

vals = counts.values
total = counts.sum()

fig, ax = plt.subplots(figsize=(6,4))
ax.pie(
    vals,
    labels=labels,
    autopct=lambda p: f"{p:.1f}\n({int(round(p*total/100))})",
    startangle=90,
    pctdistance=0.8,
    wedgeprops={"width":0.45, "edgecolor":"white"}
)
ax.set_title("Class Balance")
ax.set_aspect("equal")
plt.tight_layout()
plt.show()

```

Class Balance



```

In [31]: data = df.copy()
data["Churn"] = data["Churn"].map({"No":0, "Yes":1}).astype(int)
if "customerID" in data.columns:
    data = data.drop(columns=["customerID"])

dummies = pd.get_dummies(data, dtype=int)

y = dummies["Churn"]
X = dummies.drop(columns="Churn")
corr = X.corrwith(y)

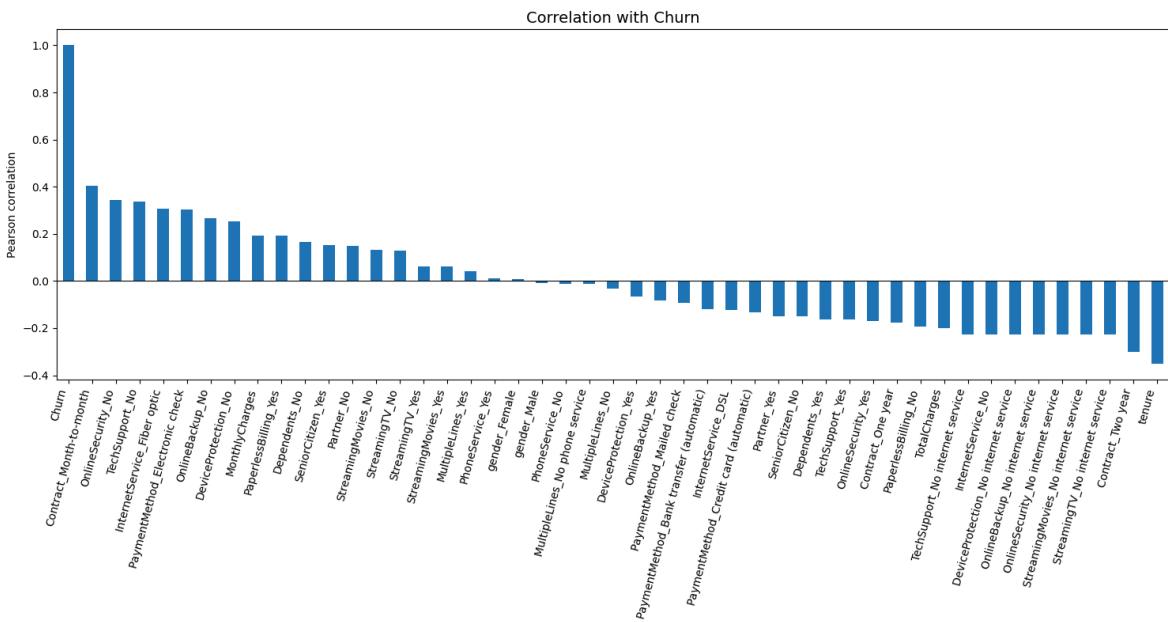
corr_with_target = pd.concat([pd.Series({"Churn": 1.0}), corr])

corr_sorted = corr_with_target.sort_values(ascending=False)

plt.figure(figsize=(15, 8))
corr_sorted.plot(kind="bar")
plt.axhline(0, color="black", linewidth=0.8)
plt.title("Correlation with Churn", fontsize=14)

```

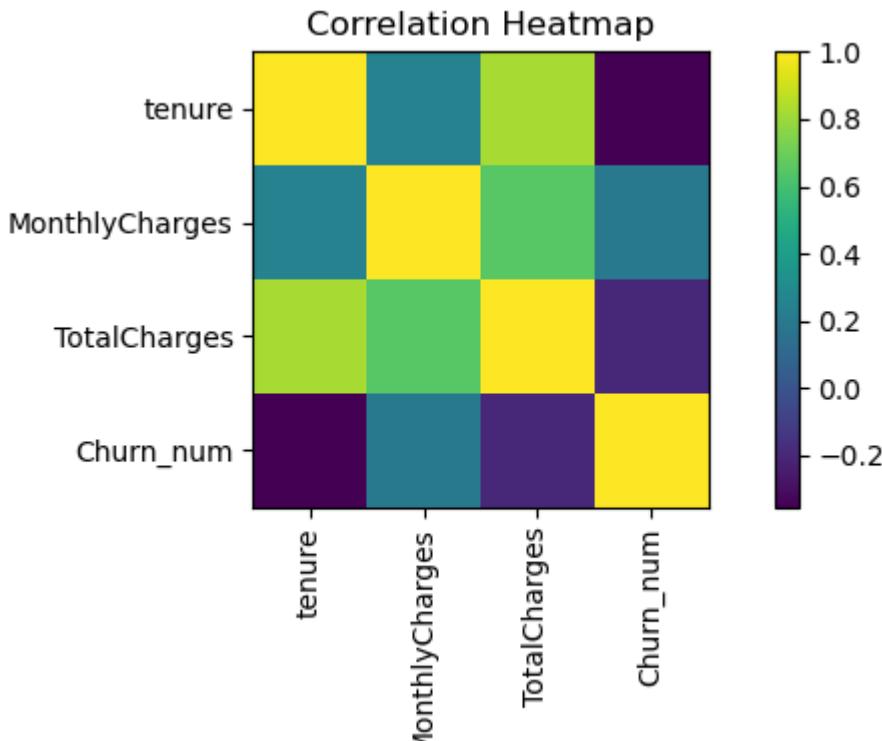
```
plt.ylabel("Pearson correlation")
plt.xticks(rotation=75, ha="right")
plt.tight_layout()
plt.show()
```



```
In [32]: corr_df = df[numerical_cols].copy()
corr_df["Churn_num"] = _churn_num
corr = corr_df.corr()

plt.imshow(corr, interpolation="nearest")
plt.title("Correlation Heatmap")
plt.colorbar()
plt.xticks(range(len(corr.columns)), corr.columns, rotation=90)
plt.yticks(range(len(corr.index)), corr.index)
plt.tight_layout(); plt.savefig("reports/figures/eda_corr_heatmap.png", dpi=200)
plt.show()

print("Top correlation with Churn_num:\n", corr["Churn_num"].drop("Churn_num").a
```



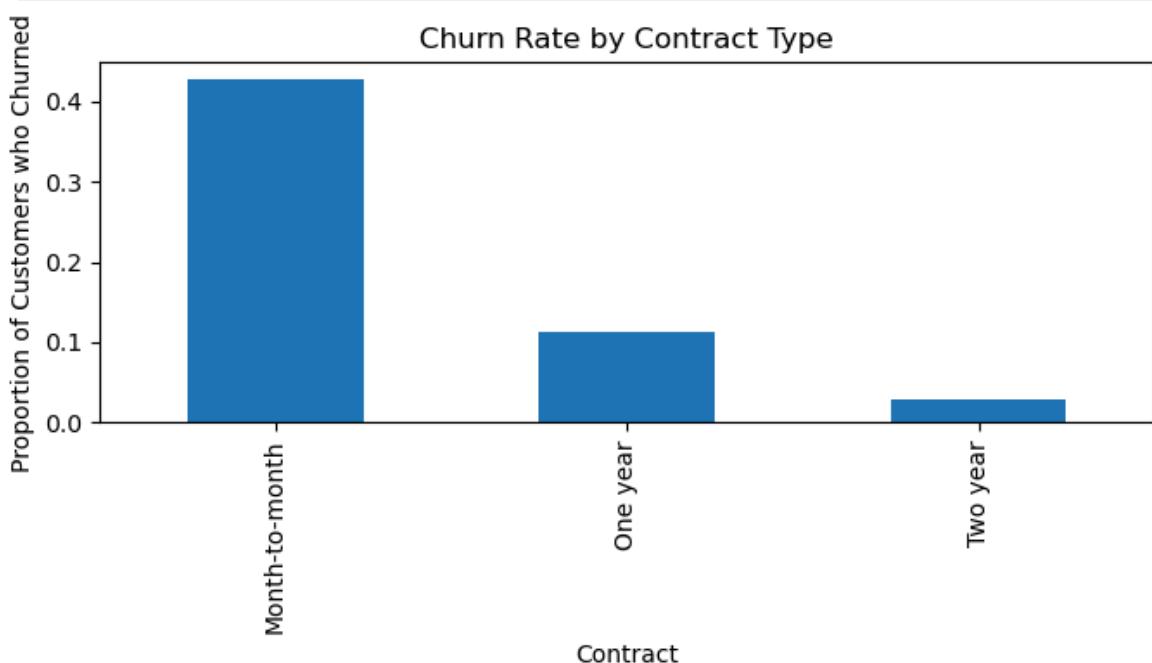
Top correlation with Churn_num:

```

tenure      0.352229
TotalCharges 0.199428
MonthlyCharges 0.193356
Name: Churn_num, dtype: float64

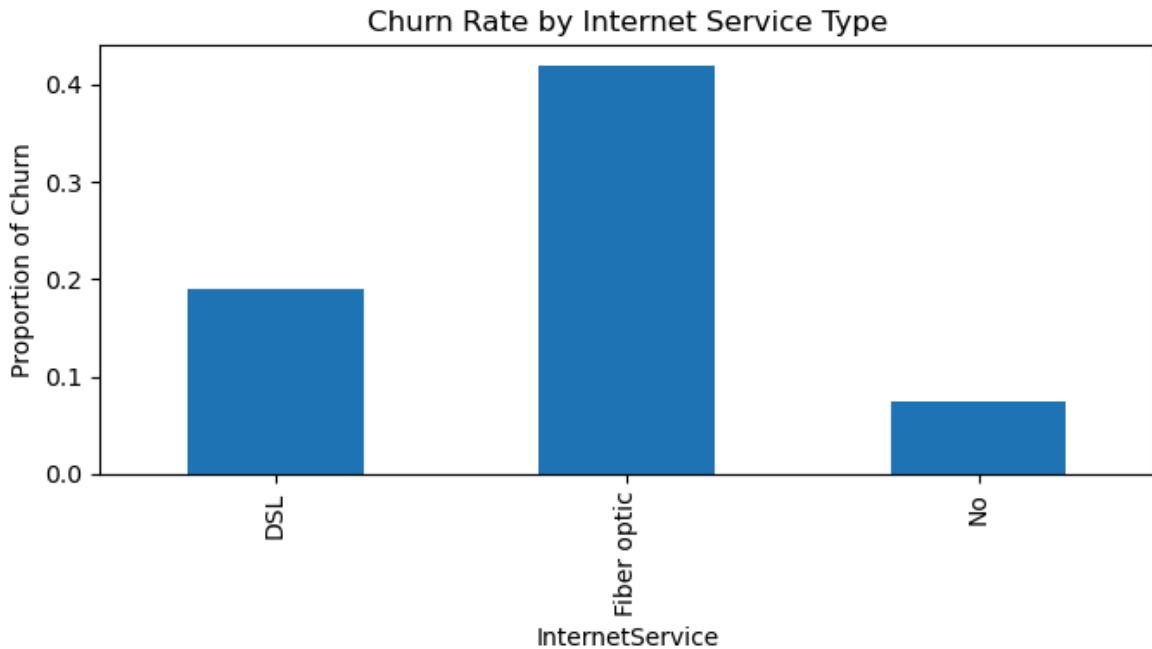
```

```
In [33]: churn_by_contract = df.groupby('Contract')['Churn'].value_counts(normalize=True)
churn_by_contract.plot(kind='bar')
plt.title('Churn Rate by Contract Type')
plt.ylabel('Proportion of Customers who Churned')
plt.tight_layout();
plt.show()
```

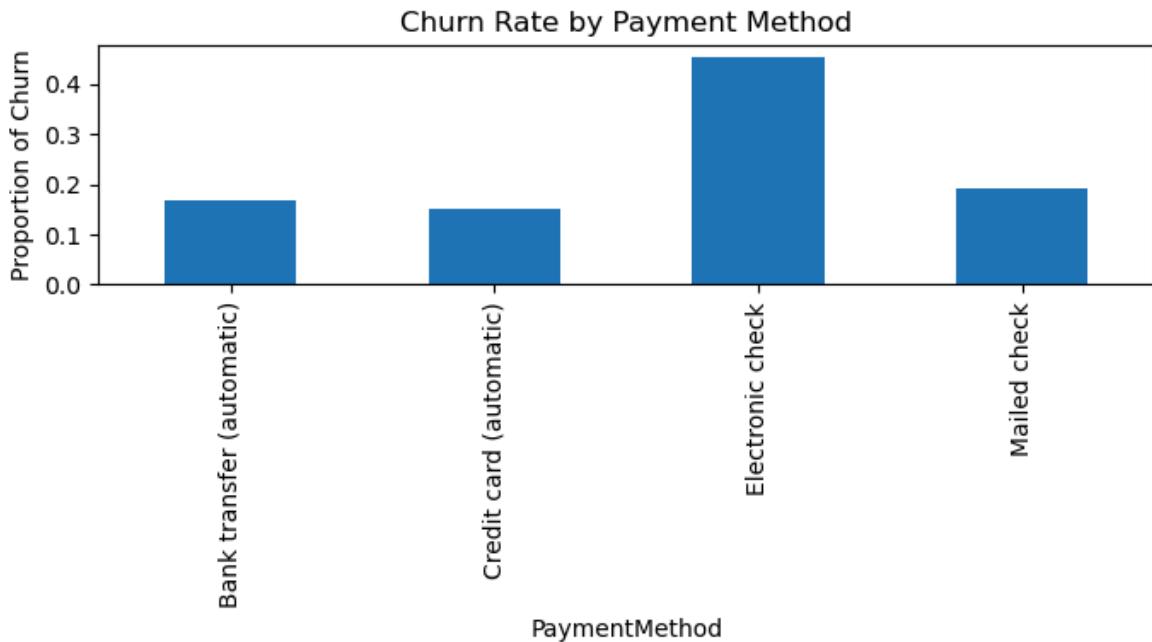


```
In [34]: churn_by_service = df.groupby('InternetService')['Churn'].value_counts(normalize=True)
churn_by_service.plot(kind='bar')
plt.title('Churn Rate by Internet Service Type')
```

```
plt.ylabel('Proportion of Churn')
plt.tight_layout(); plt.show()
```



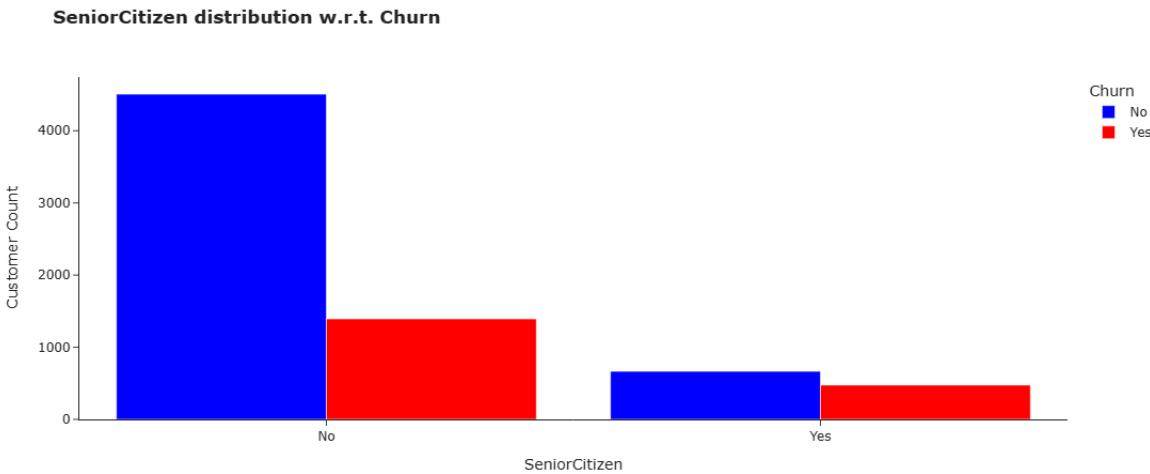
```
In [35]: churn_by_payment = df.groupby('PaymentMethod')['Churn'].value_counts(normalize=True)
churn_by_payment.plot(kind='bar')
plt.title('Churn Rate by Payment Method')
plt.ylabel('Proportion of Churn')
plt.tight_layout(); plt.show()
```



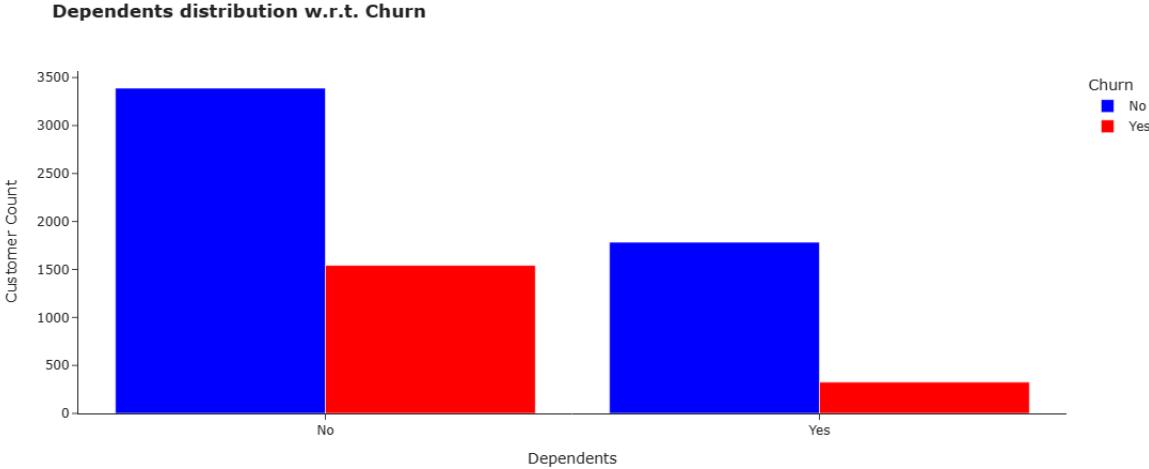
```
In [36]: top_cat = "SeniorCitizen"
color_map = {"Yes": "red", "No": "blue"}

fig = px.histogram(
    df,
    x=top_cat,
    color="Churn",
    barmode="group",
    title=f"<b>{top_cat}</b> distribution w.r.t. Churn",
    color_discrete_map=color_map)
```

```
)  
  
fig.update_layout(  
    width=700,  
    height=500,  
    bargap=0.15,  
    xaxis_title=top_cat,  
    yaxis_title="Customer Count",  
    legend_title="Churn",  
    template="simple_white"  
)  
  
fig.show()
```



```
In [37]: import plotly.express as px  
top_cat = "Dependents"  
color_map = {"Yes": "red", "No": "blue"}  
  
fig = px.histogram(  
    df,  
    x=top_cat,  
    color="Churn",  
    barmode="group",  
    title=f"<b>{top_cat}> distribution w.r.t. Churn</b>",  
    color_discrete_map=color_map  
)  
  
fig.update_layout(  
    width=700,  
    height=500,  
    bargap=0.15,  
    xaxis_title=top_cat,  
    yaxis_title="Customer Count",  
    legend_title="Churn",  
    template="simple_white"  
)  
  
fig.show()
```

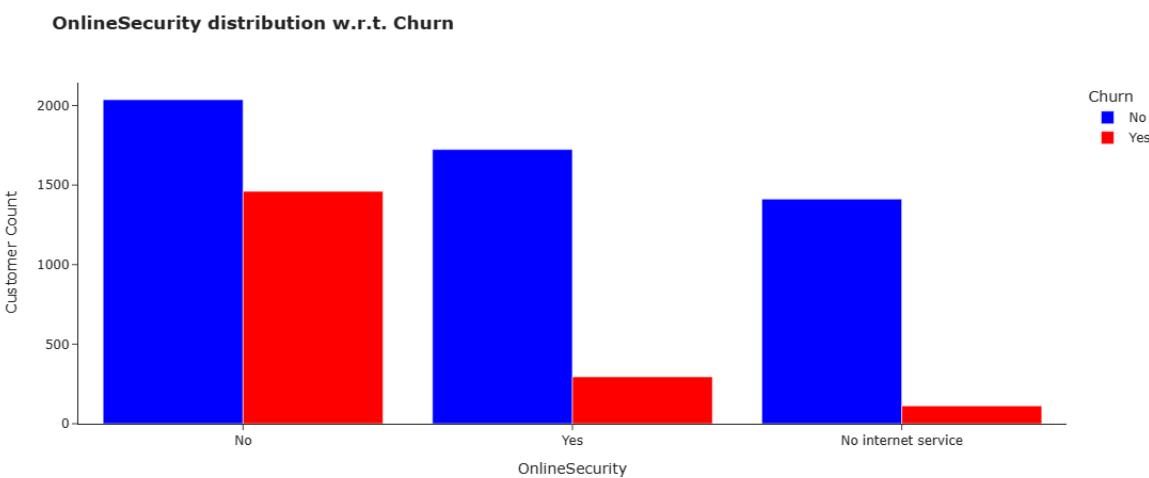


```
In [38]: top_cat = "OnlineSecurity"
color_map = {"Yes": "red", "No": "blue"}

fig = px.histogram(
    df,
    x=top_cat,
    color="Churn",
    barmode="group",
    title=f"<b>{top_cat} distribution w.r.t. Churn</b>",
    color_discrete_map=color_map
)

fig.update_layout(
    width=700,
    height=500,
    bargap=0.15,
    xaxis_title=top_cat,
    yaxis_title="Customer Count",
    legend_title="Churn",
    template="simple_white"
)

fig.show()
```



```
In [39]: top_cat = "TechSupport"
color_map = {"Yes": "red", "No": "blue"}
```

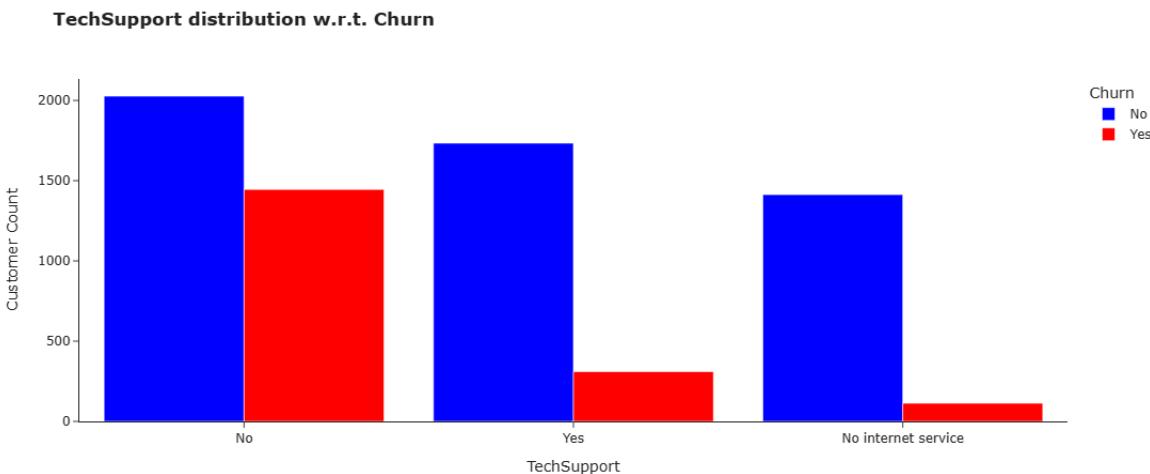
```

fig = px.histogram(
    df,
    x=top_cat,
    color="Churn",
    barmode="group",
    title=f"{top_cat} distribution w.r.t. Churn",
    color_discrete_map=color_map
)

fig.update_layout(
    width=700,
    height=500,
    bargap=0.15,
    xaxis_title=top_cat,
    yaxis_title="Customer Count",
    legend_title="Churn",
    template="simple_white"
)

fig.show()

```



Preprocessing

```

In [40]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
import matplotlib.pyplot as plt

y = df["Churn"].map({"No":0, "Yes":1}).astype(int)
X = df.drop(columns=["Churn", "customerID"], errors="ignore")

num_cols = X.select_dtypes(include=[np.number]).columns.tolist()
cat_cols = X.select_dtypes(include=["object", "category", "bool"]).columns.tolist()

X_tmp, X_test, y_tmp, y_test = train_test_split(X, y, test_size=0.20, stratify=y)
X_train, X_val, y_train, y_val = train_test_split(X_tmp, y_tmp, test_size=0.25)

print(f"Train: {X_train.shape}, Val: {X_val.shape}, Test: {X_test.shape}")

```

```

print(f"Pos rate train/val/test: {y_train.mean():.3f}/{y_val.mean():.3f}/{y_test.mean():.3f}")

numeric_pipe = Pipeline([("imputer", SimpleImputer(strategy="mean")),
                        ("scaler", StandardScaler())])
categorical_pipe = Pipeline([("imputer", SimpleImputer(strategy="most_frequent")),
                            ("onehot", OneHotEncoder(handle_unknown="ignore"))])

preprocess = ColumnTransformer([('num', numeric_pipe, num_cols),
                               ('cat', categorical_pipe, cat_cols)])

X_train_p = preprocess.fit_transform(X_train)
X_val_p = preprocess.transform(X_val)
X_test_p = preprocess.transform(X_test)

ohe = preprocess.named_transformers_["cat"].named_steps["onehot"]
ohe_names = ohe.get_feature_names_out(cat_cols)
feature_names = np.r_[num_cols, ohe_names]

print("Processed shapes:", X_train_p.shape, X_val_p.shape, X_test_p.shape)
print("Total features after preprocessing:", len(feature_names))

sizes = {"Train": len(X_train), "Val": len(X_val), "Test": len(X_test)}
plt.figure(figsize=(6,4))
plt.bar(list(sizes.keys()), list(sizes.values()))
plt.title("Dataset Split Counts")
plt.ylabel("Rows")
plt.tight_layout()
plt.show()

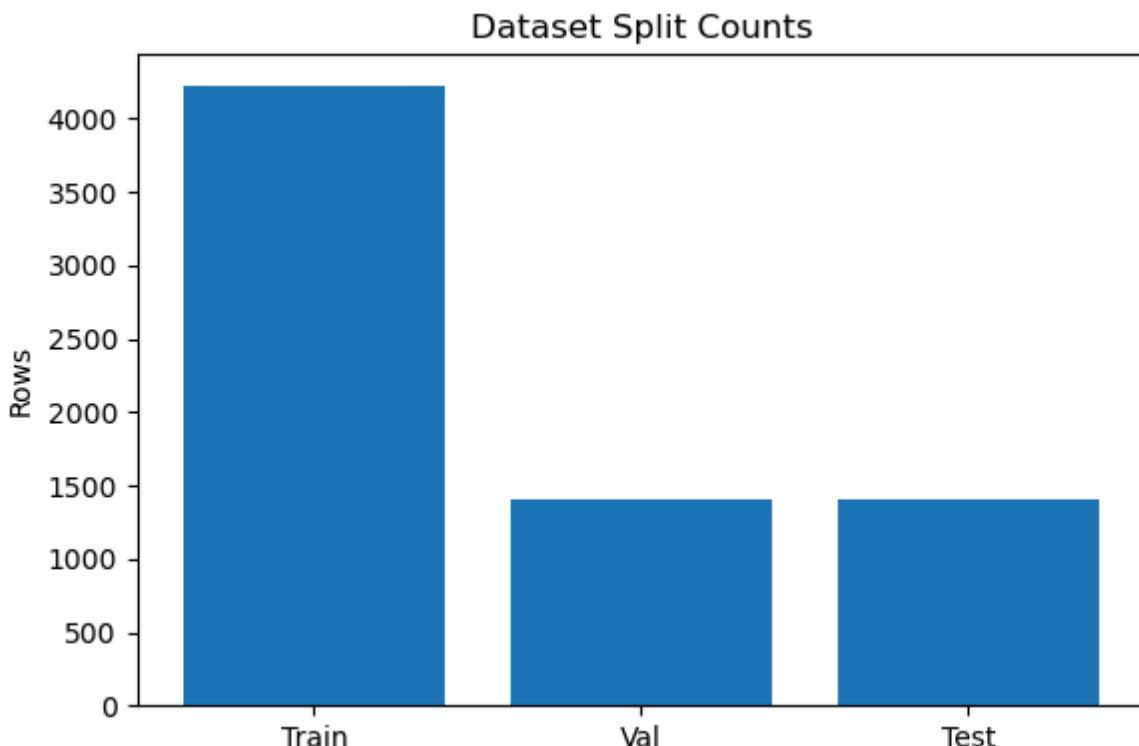
```

Train: (4225, 19), Val: (1409, 19), Test: (1409, 19)

Pos rate train/val/test: 0.265/0.265/0.265

Processed shapes: (4225, 46) (1409, 46) (1409, 46)

Total features after preprocessing: 46



In [58]:

```

import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split, cross_val_score

```

```
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
import xgboost as xgb
from sklearn.metrics import (accuracy_score, precision_score, recall_score,
                             f1_score, roc_auc_score, confusion_matrix,
                             classification_report)
from sklearn.inspection import permutation_importance
import shap
import matplotlib.pyplot as plt
import seaborn as sns

# Model Training

# Baseline: predict majority class
baseline_accuracy = 1 - y_train.mean() # Assuming no churn is majority
print(f"BASELINE ACCURACY (predict majority class): {baseline_accuracy:.4f}")

models = {
    "Logistic Regression": LogisticRegression(max_iter=1000, random_state=123),
    "Decision Tree": DecisionTreeClassifier(random_state=123),
    "Random Forest": RandomForestClassifier(n_estimators=100, random_state=123),
    "Gradient Boosting": GradientBoostingClassifier(n_estimators=100, random_state=123)
}

param_grids = {
    "Logistic Regression": [
        {"C": 0.01, "solver": "lbfgs"},
        {"C": 0.1, "solver": "lbfgs"},
        {"C": 1.0, "solver": "lbfgs"},
        {"C": 10.0, "solver": "lbfgs"}
    ],
    "Decision Tree": [
        {"max_depth": 5, "min_samples_split": 20},
        {"max_depth": 10, "min_samples_split": 20},
        {"max_depth": 15, "min_samples_split": 10},
        {"max_depth": None, "min_samples_split": 10}
    ],
    "Random Forest": [
        {"n_estimators": 50, "max_depth": 10, "min_samples_split": 20},
        {"n_estimators": 100, "max_depth": 15, "min_samples_split": 10},
        {"n_estimators": 200, "max_depth": 20, "min_samples_split": 5}
    ],
    "Gradient Boosting": [
        {"n_estimators": 50, "learning_rate": 0.1, "max_depth": 3},
        {"n_estimators": 100, "learning_rate": 0.1, "max_depth": 3},
        {"n_estimators": 100, "learning_rate": 0.05, "max_depth": 5}
    ]
}

# Training and Evaluation

results = []
best_models = {}
```

```

for model_name, base_model in models.items():
    print(f"\nTraining {model_name}...")

    best_score = 0
    best_params = None
    best_model_instance = None

    for params in param_grids[model_name]:
        model = base_model.__class__(**params, random_state=123)
        model.fit(X_train_p, y_train)
        val_score = roc_auc_score(y_val, model.predict_proba(X_val_p)[:, 1])

        if val_score > best_score:
            best_score = val_score
            best_params = params
            best_model_instance = model

    cv_scores = cross_val_score(
        best_model_instance, X_train_p, y_train,
        cv=5, scoring='roc_auc', n_jobs=-1
    )

    test_scores = []
    if model_name in ["Random Forest", "Gradient Boosting"]:
        for seed in range(123, 128):
            temp_model = best_model_instance.__class__(**best_params, random_state=seed)
            temp_model.fit(X_train_p, y_train)
            test_pred_proba = temp_model.predict_proba(X_test_p)[:, 1]
            test_scores.append(roc_auc_score(y_test, test_pred_proba))
    else:
        test_pred_proba = best_model_instance.predict_proba(X_test_p)[:, 1]
        test_scores = [roc_auc_score(y_test, test_pred_proba)]

    y_test_pred = best_model_instance.predict(X_test_p)
    y_test_proba = best_model_instance.predict_proba(X_test_p)[:, 1]

    test_accuracy = accuracy_score(y_test, y_test_pred)
    test_precision = precision_score(y_test, y_test_pred)
    test_recall = recall_score(y_test, y_test_pred)
    test_f1 = f1_score(y_test, y_test_pred)
    test_auc = roc_auc_score(y_test, y_test_proba)

    results.append({
        "Model": model_name,
        "Best Params": str(best_params),
        "CV AUC": f"{cv_scores.mean():.4f}±{cv_scores.std():.4f}",
        "Test Accuracy": test_accuracy,
        "Test Precision": test_precision,
        "Test Recall": test_recall,
        "Test F1": test_f1,
        "Test AUC": test_auc,
        "Test AUC(mean, std)": f"{np.mean(test_scores):.4f}±{np.std(test_scores):.4f}"
    })

    best_models[model_name] = best_model_instance

    print(f"  Best params: {best_params}")
    print(f"  CV AUC: {cv_scores.mean():.4f} ± {cv_scores.std():.4f}")
    print(f"  Test AUC: {np.mean(test_scores):.4f} ± {np.std(test_scores):.4f}")

```

```

# Performance Comparison Table

results_df = pd.DataFrame(results)

print("MODEL PERFORMANCE SUMMARY")
print(results_df.to_string(index=False))

# Find best model
best_model_name = results_df.loc[results_df["Test AUC"].idxmax(), "Model"]
best_model = best_models[best_model_name]
print(f"Best Model: {best_model_name} (Test AUC: {results_df['Test AUC'].max():.2f})")

best_accuracy = results_df.loc[results_df["Test AUC"].idxmax(), "Test Accuracy"]
std_above_baseline = (best_accuracy - baseline_accuracy) / (baseline_accuracy * std)
print(f"Best model is {std_above_baseline:.2f} standard deviations above baseline accuracy")

# Visualization

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# AUC comparison
ax1 = axes[0]
models_list = results_df["Model"].tolist()
auc_scores = results_df["Test AUC"].tolist()
colors = plt.cm.viridis(np.linspace(0, 1, len(models_list)))

bars = ax1.banh(models_list, auc_scores, color=colors)
ax1.axvline(x=0.5, color='red', linestyle='--', label='Random Guess', alpha=0.7)
ax1.set_xlabel('Test AUC Score', fontsize=12)
ax1.set_title('Model Performance Comparison (AUC)', fontsize=14, fontweight='bold')
ax1.legend()
ax1.grid(axis='x', alpha=0.3)
for i, (bar, score) in enumerate(zip(bars, auc_scores)):
    ax1.text(score + 0.01, i, f'{score:.3f}', va='center', fontsize=10)

# Multiple metrics for best model
ax2 = axes[1]
best_idx = results_df["Test AUC"].idxmax()
metrics = ['Accuracy', 'Precision', 'Recall', 'F1', 'AUC']
values = [
    results_df.loc[best_idx, 'Test Accuracy'],
    results_df.loc[best_idx, 'Test Precision'],
    results_df.loc[best_idx, 'Test Recall'],
    results_df.loc[best_idx, 'Test F1'],
    results_df.loc[best_idx, 'Test AUC']
]
colors2 = plt.cm.plasma(np.linspace(0, 1, len(metrics)))
bars2 = ax2.bar(metrics, values, color=colors2)
ax2.set_ylabel('Score', fontsize=12)
ax2.set_title(f'{best_model_name} - All Metrics', fontsize=14, fontweight='bold')
ax2.set_ylim([0, 1])
ax2.grid(axis='y', alpha=0.3)
for bar, val in zip(bars2, values):
    height = bar.get_height()
    ax2.text(bar.get_x() + bar.get_width()/2., height + 0.02,
             f'{val:.3f}', ha='center', va='bottom', fontsize=10)

```

```
plt.tight_layout()
plt.show()

# Feature Importance (3+ Methods)

print("FEATURE IMPORTANCE ANALYSIS")

# Built-in feature importance
if hasattr(best_model, 'feature_importances_'):
    importance_builtin = best_model.feature_importances_
    feature_imp_df = pd.DataFrame({
        'Feature': feature_names,
        'Importance': importance_builtin
    }).sort_values('Importance', ascending=False)

    print("Top 10 Features (Built-in Importance):")
    print(feature_imp_df.head(10).to_string(index=False))
    print()

    plt.figure(figsize=(10, 6))
    top_n = 15
    top_features = feature_imp_df.head(top_n)
    plt.barh(range(top_n), top_features['Importance'].values)
    plt.yticks(range(top_n), top_features['Feature'].values)
    plt.xlabel('Feature Importance', fontsize=12)
    plt.title(f'Top {top_n} Features - Built-in Importance ({best_model_name})',
              fontsize=14, fontweight='bold')
    plt.gca().invert_yaxis()
    plt.tight_layout()
    plt.show()

# Permutation Importance
print("\n Permutation Importance")
perm_importance = permutation_importance(
    best_model, X_test_p, y_test,
    n_repeats=10, random_state=123, n_jobs=-1
)

perm_imp_df = pd.DataFrame({
    'Feature': feature_names,
    'Importance': perm_importance.importances_mean,
    'Std': perm_importance.importances_std
}).sort_values('Importance', ascending=False)

print("\nTop 10 Features (Permutation Importance):")
print(perm_imp_df.head(10).to_string(index=False))

plt.figure(figsize=(10, 6))
top_n = 15
top_features_perm = perm_imp_df.head(top_n)
y_pos = range(top_n)
plt.barh(y_pos, top_features_perm['Importance'].values,
         xerr=top_features_perm['Std'].values, alpha=0.7)
plt.yticks(y_pos, top_features_perm['Feature'].values)
plt.xlabel('Permutation Importance', fontsize=12)
plt.title(f'Top {top_n} Features - Permutation Importance',
          fontsize=14, fontweight='bold')
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()
```

```

# Coefficients for Logistic Regression or alternative method
if "Logistic" in best_model_name:
    coef_importance = np.abs(best_models["Logistic Regression"].coef_[0])
    coef_imp_df = pd.DataFrame({
        'Feature': feature_names,
        'Importance': coef_importance
    }).sort_values('Importance', ascending=False)

    print("\nTop 10 Features (Coefficient Magnitude - Logistic Regression):")
    print(coef_imp_df.head(10).to_string(index=False))

    plt.figure(figsize=(10, 6))
    top_n = 15
    top_features_coef = coef_imp_df.head(top_n)
    plt.barh(range(top_n), top_features_coef['Importance'].values, color='coral')
    plt.yticks(range(top_n), top_features_coef['Feature'].values)
    plt.xlabel('Absolute Coefficient Value', fontsize=12)
    plt.title(f'Top {top_n} Features - Coefficient Magnitude',
              fontsize=14, fontweight='bold')
    plt.gca().invert_yaxis()
    plt.tight_layout()
    plt.show()

# SHAP Values for Local Interpretability

print("SHAP ANALYSIS")

if best_model_name in ["Random Forest", "Gradient Boosting", "Decision Tree"]:
    X_shap = X_test_p
    explainer = shap.TreeExplainer(best_model)
    shap_values = explainer.shap_values(X_shap)
else:
    X_shap = X_test_p[:100]
    background = shap.sample(X_train_p, 100)
    explainer = shap.KernelExplainer(best_model.predict_proba, background)
    shap_values = explainer.shap_values(X_shap)

shap_values_arr = np.array(shap_values)

if isinstance(shap_values, list):
    shap_values_plot = shap_values[1]
    base_val = explainer.expected_value[1]
else:
    if shap_values_arr.ndim == 3:
        shap_values_plot = shap_values_arr[:, :, 1]
        if hasattr(explainer.expected_value, "__len__"):
            base_val = explainer.expected_value[1]
        else:
            base_val = explainer.expected_value
    else:
        shap_values_plot = shap_values_arr
        if hasattr(explainer.expected_value, "__len__"):
            base_val = explainer.expected_value[1]
        else:
            base_val = explainer.expected_value

if hasattr(X_shap, "toarray"):
    X_shap_plot = X_shap.toarray()
else:

```

```

X_shap_plot = X_shap

plt.figure(figsize=(10, 8))
shap.summary_plot(
    shap_values_plot,
    X_shap_plot,
    feature_names=feature_names,
    show=False
)
plt.title(
    'SHAP Summary Plot - Feature Impact on Predictions',
    fontsize=14,
    fontweight='bold',
    pad=20
)
plt.tight_layout()
plt.show()

# SHAP Bar Plot
plt.figure(figsize=(10, 6))
shap.summary_plot(
    shap_values_plot,
    X_shap_plot,
    feature_names=feature_names,
    plot_type="bar",
    show=False
)
plt.title(
    'SHAP Feature Importance - Mean Absolute Impact',
    fontsize=14,
    fontweight='bold'
)
plt.tight_layout()
plt.show()

# waterfall plot for a single instance

sample_idx = 0

sample_data = X_shap[sample_idx]
if hasattr(sample_data, "toarray"):
    sample_data = sample_data.toarray().ravel()

values_1d = shap_values_plot[sample_idx]

shap_expl = shap.Explanation(
    values=values_1d,
    base_values=base_val,
    data=sample_data,
    feature_names=feature_names
)

shap.plots.waterfall(shap_expl, max_display=15)

print("\nKey Findings:")
print(f"1. Best Model: {best_model_name}")
print(f"2. Test AUC: {results_df['Test AUC'].max():.4f}")
print(f"3. Performance is {std_above_baseline:.2f} std above baseline")
print(f"4. Most important features identified through multiple methods")
print(f"5. SHAP analysis provides local explanations for individual predictions"

```

BASELINE ACCURACY (predict majority class): 0.7347

Training Logistic Regression...

Best params: {'C': 10.0, 'solver': 'lbfgs'}
 CV AUC: 0.8406 ± 0.0134
 Test AUC: 0.8437 ± 0.0000

Training Decision Tree...

Best params: {'max_depth': 5, 'min_samples_split': 20}
 CV AUC: 0.8214 ± 0.0097
 Test AUC: 0.8183 ± 0.0000

Training Random Forest...

Best params: {'n_estimators': 50, 'max_depth': 10, 'min_samples_split': 20}
 CV AUC: 0.8415 ± 0.0157
 Test AUC: 0.8390 ± 0.0011

Training Gradient Boosting...

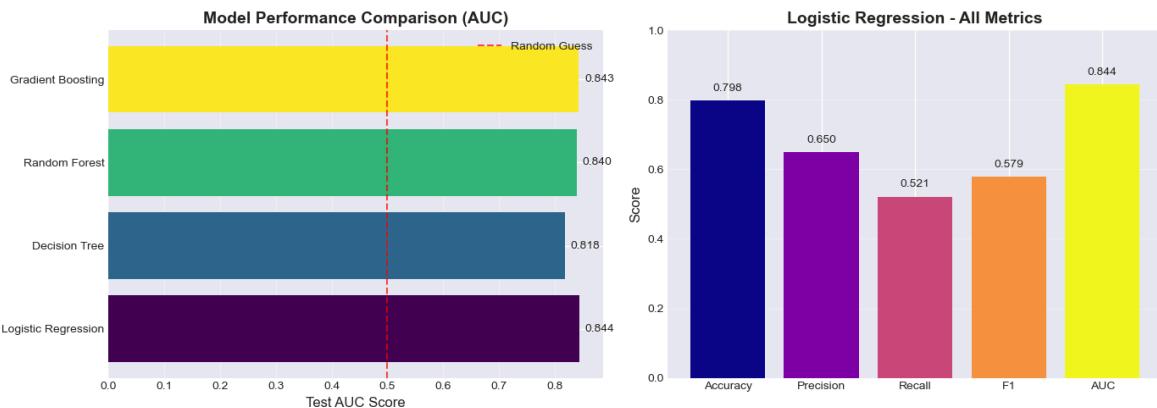
Best params: {'n_estimators': 50, 'learning_rate': 0.1, 'max_depth': 3}
 CV AUC: 0.8433 ± 0.0142
 Test AUC: 0.8428 ± 0.0000

MODEL PERFORMANCE SUMMARY

	Model	Best Param					
s	CV AUC	Test Accuracy	Test Precision	Test Recall	Test F1	Test AUC	Te
st AUC(mean, std)							
Logistic Regression							
s'}	0.8406±0.0134	0.798439	0.650000	0.521390	0.578635	0.843724	
0.8437±0.0000							
Decision Tree							
0}	0.8214±0.0097	0.791341	0.613636	0.577540	0.595041	0.818321	
0.8183±0.0000							
Random Forest	{'n_estimators': 50, 'max_depth': 10, 'min_samples_split': 2						
0}	0.8415±0.0157	0.798439	0.659574	0.497326	0.567073	0.839583	
0.8390±0.0011							
Gradient Boosting							
3}	0.8433±0.0142	0.801987	0.664360	0.513369	0.579186	0.842808	
0.8428±0.0000							

Best Model: Logistic Regression (Test AUC: 0.8437)

Best model is 0.14 standard deviations above baseline



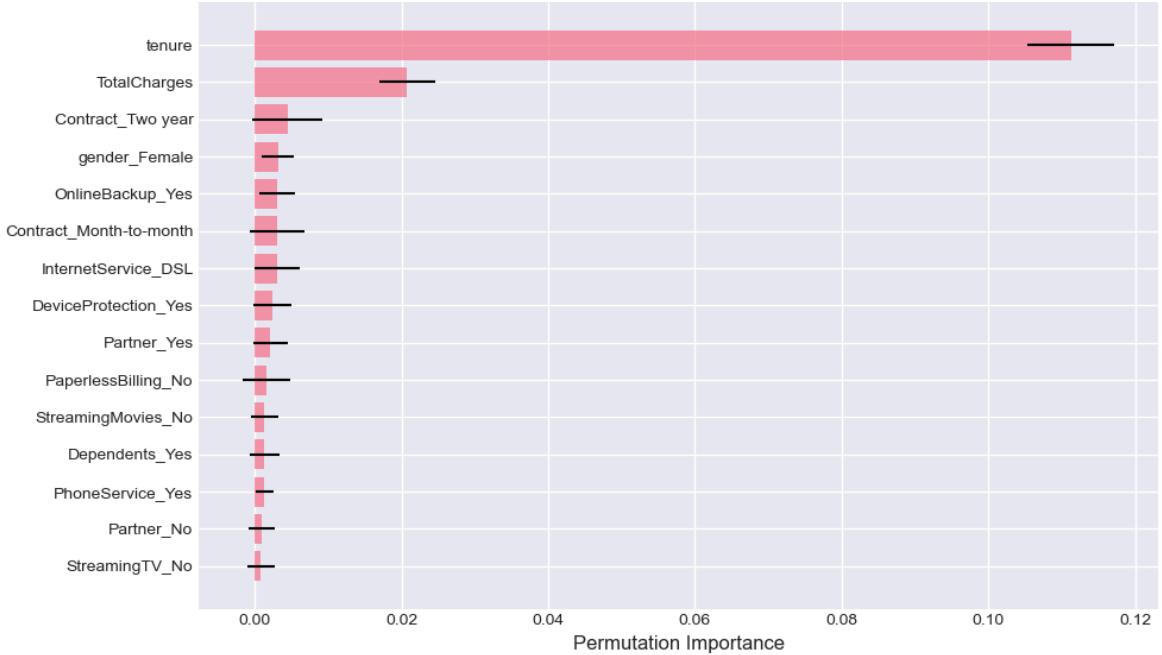
FEATURE IMPORTANCE ANALYSIS

Permutation Importance

Top 10 Features (Permutation Importance):

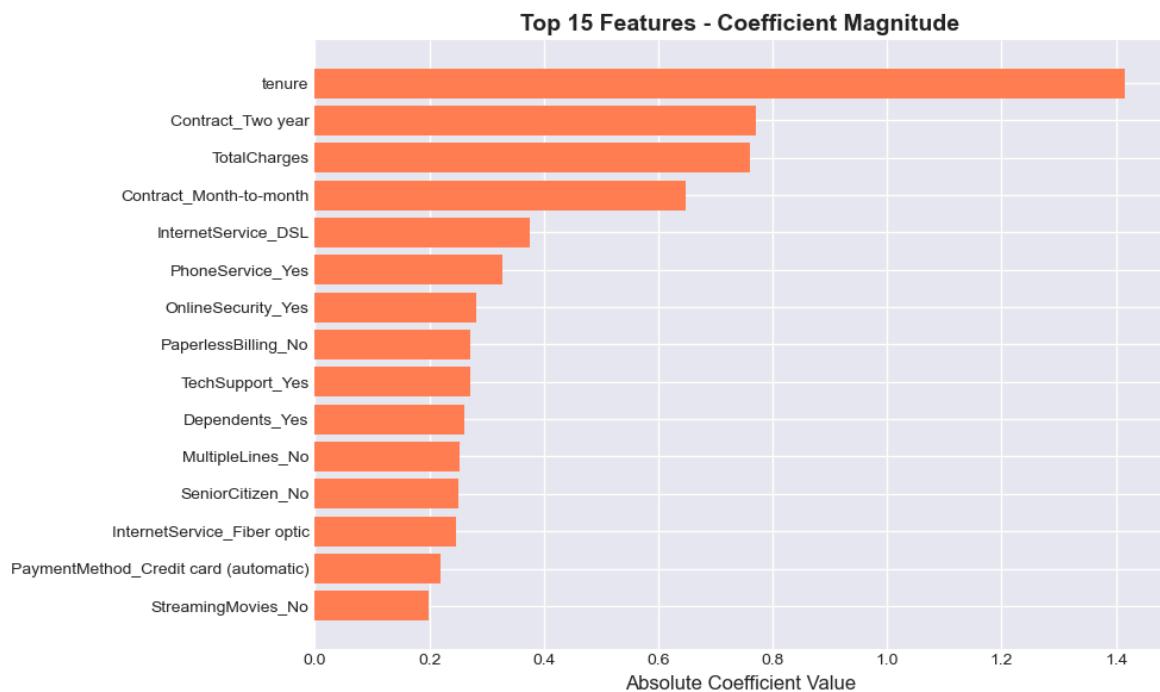
	Feature	Importance	Std
	tenure	0.111285	0.005902
	TotalCharges	0.020795	0.003743
	Contract_Two year	0.004471	0.004761
	gender_Female	0.003194	0.002205
	OnlineBackup_Yes	0.003052	0.002418
	Contract_Month-to-month	0.003052	0.003675
	InternetService_DSL	0.002981	0.003074
	DeviceProtection_Yes	0.002342	0.002580
	Partner_Yes	0.002129	0.002311
	PaperlessBilling_No	0.001561	0.003202

Top 15 Features - Permutation Importance



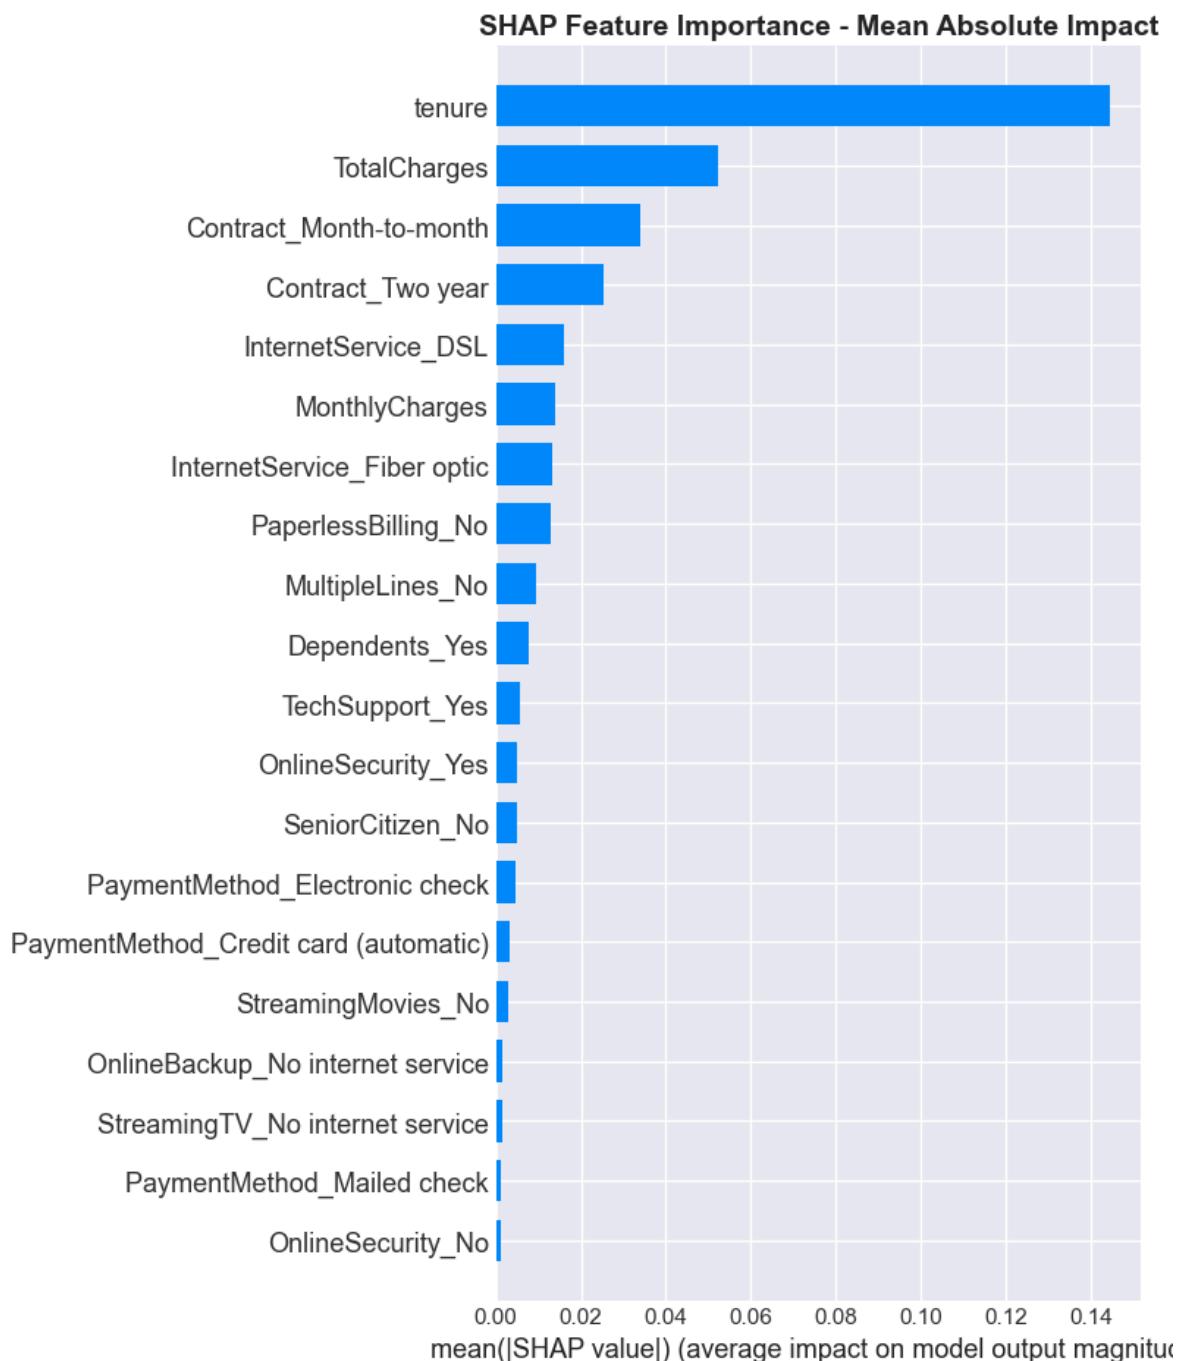
Top 10 Features (Coefficient Magnitude - Logistic Regression):

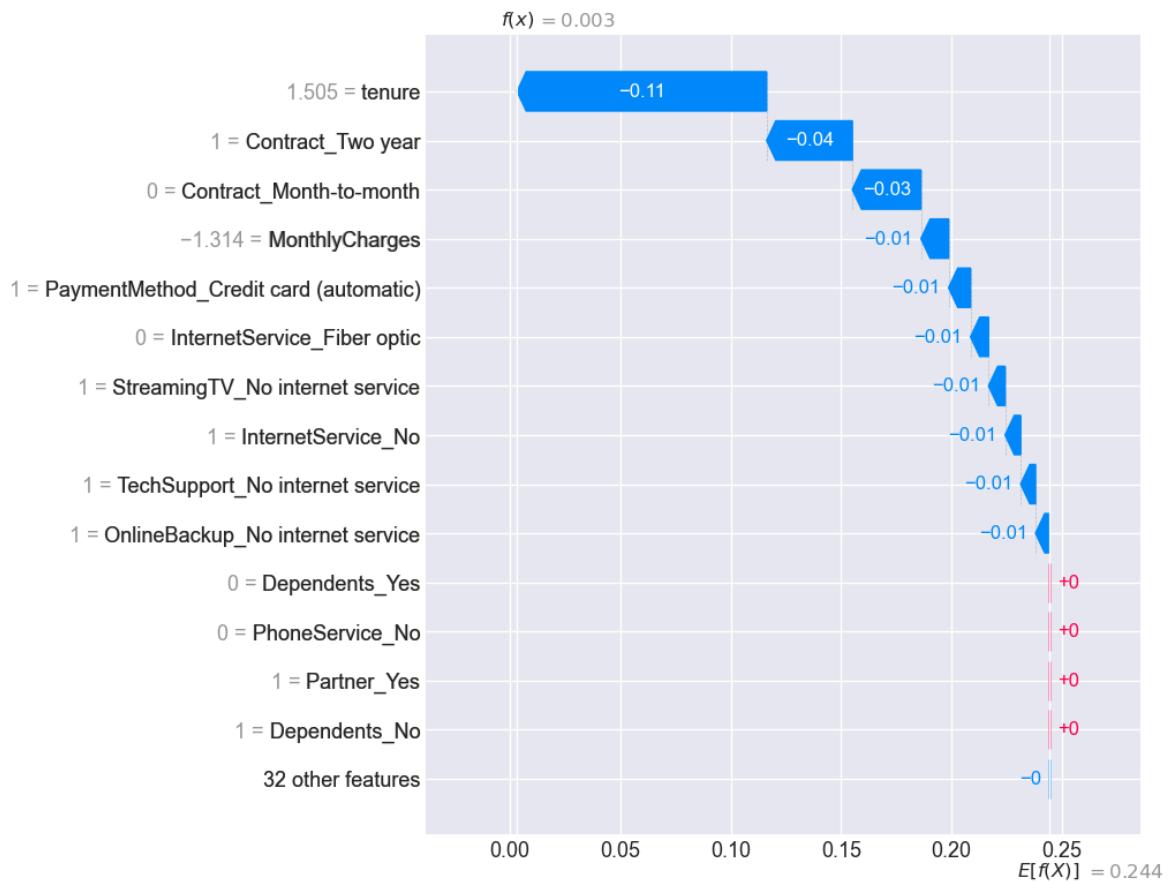
	Feature	Importance
	tenure	1.415206
	Contract_Two year	0.771120
	TotalCharges	0.761122
	Contract_Month-to-month	0.647923
	InternetService_DSL	0.374971
	PhoneService_Yes	0.327445
	OnlineSecurity_Yes	0.281240
	PaperlessBilling_No	0.272287
	TechSupport_Yes	0.270703
	Dependents_Yes	0.261183

**SHAP ANALYSIS**

0% | 0/100 [00:00<?, ?it/s]

SHAP Summary Plot - Feature Impact on Predictions





Key Findings:

1. Best Model: Logistic Regression
2. Test AUC: 0.8437
3. Performance is 0.14 std above baseline
4. Most important features identified through multiple methods
5. SHAP analysis provides local explanations for individual predictions

In []: